

Copyright 2018 The TensorFlow Authors.

In [1]:

```
#@title Licensed under the Apache License, Version 2.0 (the "License");
# you may not use this file except in compliance with the License.
# You may obtain a copy of the License at
#
# https://www.apache.org/licenses/LICENSE-2.0
#
# Unless required by applicable law or agreed to in writing, software
# distributed under the License is distributed on an "AS IS" BASIS,
# WITHOUT WARRANTIES OR CONDITIONS OF ANY KIND, either express or implied.
# See the License for the specific language governing permissions and
# limitations under the License.
```

In [2]:

```
#@title MIT License
#
# Copyright (c) 2017 François Chollet
#
# Permission is hereby granted, free of charge, to any person obtaining a
# copy of this software and associated documentation files (the "Software"),
# to deal in the Software without restriction, including without limitation
# the rights to use, copy, modify, merge, publish, distribute, sublicense,
# and/or sell copies of the Software, and to permit persons to whom the
# Software is furnished to do so, subject to the following conditions:
#
# The above copyright notice and this permission notice shall be included in
# all copies or substantial portions of the Software.
#
# THE SOFTWARE IS PROVIDED "AS IS", WITHOUT WARRANTY OF ANY KIND, EXPRESS OR
# IMPLIED, INCLUDING BUT NOT LIMITED TO THE WARRANTIES OF MERCHANTABILITY,
# FITNESS FOR A PARTICULAR PURPOSE AND NONINFRINGEMENT. IN NO EVENT SHALL
# THE AUTHORS OR COPYRIGHT HOLDERS BE LIABLE FOR ANY CLAIM, DAMAGES OR OTHER
# LIABILITY, WHETHER IN AN ACTION OF CONTRACT, TORT OR OTHERWISE, ARISING
# FROM, OUT OF OR IN CONNECTION WITH THE SOFTWARE OR THE USE OR OTHER
# DEALINGS IN THE SOFTWARE.
```

Text classification with movie reviews



[View on TensorFlow.org](https://www.tensorflow.org/tutorials/keras/basic_text_classification)

https://www.tensorflow.org/tutorials/keras/basic_text_classification



[Run in Google](https://colab.research.google.com/github/tensorflow/docs/blob/master/site/en/tutorials/keras/basic_text_classification)

https://colab.research.google.com/github/tensorflow/docs/blob/master/site/en/tutorials/keras/basic_text_classification

This notebook classifies movie reviews as *positive* or *negative* using the text of the review. This is an example of *binary*—or two-class—classification, an important and widely applicable kind of machine learning problem.

We'll use the [IMDB dataset](https://www.tensorflow.org/api_docs/python/tf/keras/datasets/imdb) (https://www.tensorflow.org/api_docs/python/tf/keras/datasets/imdb) that contains the text of 50,000 movie reviews from the [Internet Movie Database](https://www.imdb.com/) (<https://www.imdb.com/>). These are split into 25,000 reviews for training and 25,000 reviews for testing. The training and testing sets are *balanced*, meaning they contain an equal number of positive and negative reviews.

This notebook uses [tf.keras](https://www.tensorflow.org/guide/keras) (<https://www.tensorflow.org/guide/keras>), a high-level API to build and train models in TensorFlow. For a more advanced text classification tutorial using `tf.keras`, see the [MLCC Text Classification Guide](https://developers.google.com/machine-learning/guides/text-classification/) (<https://developers.google.com/machine-learning/guides/text-classification/>).

In [3]:

```
# keras.datasets.imdb is broken in 1.13 and 1.14, by np 1.16.3  
!pip install tf_nightly
```

```
Collecting tf_nightly
  Downloading https://files.pythonhosted.org/packages/3f/4d/29da385b6ebcf100419d597d8caabd1c8b1cb5670529f8ab
15d2829e0cbe/tf_nightly-1.15.0.dev20190821-cp36-cp36m-manylinux2010_x86_64.whl (110.8MB)
|████████████████████████████████████████| 110.8MB 96.9MB/s eta 0:00:01 |████████████████████████████████████████| 2
4.2MB 5.4MB/s eta 0:00:17 |████████████████████████████████████████| 74.4MB 96.9MB/s eta 0:00:01
Requirement already satisfied: astor>=0.6.0 in /usr/local/lib/python3.6/dist-packages (from tf_nightly) (0.
8.0)
Requirement already satisfied: wheel>=0.26 in /usr/lib/python3/dist-packages (from tf_nightly) (0.30.0)
Collecting opt-einsum>=2.3.2 (from tf_nightly)
  Downloading https://files.pythonhosted.org/packages/c0/1a/ab5683d8e450e380052d3a3e77bb2c9dffa878058f583587
c3875041fb63/opt_einsum-3.0.1.tar.gz (66kB)
|████████████████████████████████████████| 71kB 38.4MB/s eta 0:00:01
Collecting tf-estimator-nightly (from tf_nightly)
  Downloading https://files.pythonhosted.org/packages/9c/bc/4aea89a134fdf4e1109951569c7fee4119e20aeac39253ba
2ec15d6181b5/tf_estimator_nightly-2.0.0-py2.py3-none-any.whl (450kB)
|████████████████████████████████████████| 450kB 30.4MB/s eta 0:00:01
Requirement already satisfied: google-pasta>=0.1.6 in /usr/local/lib/python3.6/dist-packages (from tf_nightl
y) (0.1.7)
Requirement already satisfied: six>=1.10.0 in /usr/lib/python3/dist-packages (from tf_nightly) (1.11.0)
Requirement already satisfied: keras-applications>=1.0.8 in /usr/local/lib/python3.6/dist-packages (from tf_
nightly) (1.0.8)
Requirement already satisfied: keras-preprocessing>=1.0.5 in /usr/local/lib/python3.6/dist-packages (from tf
_nightly) (1.1.0)
Collecting tb-nightly<1.16.0a0,>=1.15.0a0 (from tf_nightly)
  Downloading https://files.pythonhosted.org/packages/cd/67/301f684e269786d65296d97ed8cae65d06864c336de8beb5
3f92bf84fb82/tb_nightly-1.15.0a20190911-py3-none-any.whl (3.8MB)
|████████████████████████████████████████| 3.8MB 38.4MB/s eta 0:00:01
Requirement already satisfied: numpy<2.0,>=1.16.0 in /usr/local/lib/python3.6/dist-packages (from tf_nightl
y) (1.17.2)
Requirement already satisfied: grpcio>=1.8.6 in /usr/local/lib/python3.6/dist-packages (from tf_nightly) (1.
23.0)
Requirement already satisfied: absl-py>=0.7.0 in /usr/local/lib/python3.6/dist-packages (from tf_nightly)
(0.8.0)
Requirement already satisfied: wrapt>=1.11.1 in /usr/local/lib/python3.6/dist-packages (from tf_nightly) (1.
11.2)
Requirement already satisfied: protobuf>=3.6.1 in /usr/local/lib/python3.6/dist-packages (from tf_nightly)
(3.9.1)
Requirement already satisfied: termcolor>=1.1.0 in /usr/local/lib/python3.6/dist-packages (from tf_nightly)
(1.1.0)
Requirement already satisfied: gast>=0.2.0 in /usr/local/lib/python3.6/dist-packages (from tf_nightly) (0.3.
0)
Requirement already satisfied: h5py in /usr/local/lib/python3.6/dist-packages (from keras-applications>=1.0.
8->tf_nightly) (2.9.0)
Requirement already satisfied: werkzeug>=0.11.15 in /usr/local/lib/python3.6/dist-packages (from tb-nightly<
1.16.0a0,>=1.15.0a0->tf_nightly) (0.15.6)
Requirement already satisfied: setuptools>=41.0.0 in /usr/local/lib/python3.6/dist-packages (from tb-nightly
<1.16.0a0,>=1.15.0a0->tf_nightly) (41.0.1)
```

```
Requirement already satisfied: markdown>=2.6.8 in /usr/local/lib/python3.6/dist-packages (from tb-nightly<1.16.0a0,>=1.15.0a0->tf-nightly) (3.1.1)
Building wheels for collected packages: opt-einsum
  Building wheel for opt-einsum (setup.py) ... done
  Stored in directory: /root/.cache/pip/wheels/91/98/8d/10e3d4e04c959597a411b91acd3695e9e2d210e68ce3427aad
Successfully built opt-einsum
Installing collected packages: opt-einsum, tf-estimator-nightly, tb-nightly, tf-nightly
Successfully installed opt-einsum-3.0.1 tb-nightly-1.15.0a20190911 tf-estimator-nightly-2.0.0 tf-nightly-1.15.0.dev20190821
WARNING: You are using pip version 19.1.1, however version 19.2.3 is available.
You should consider upgrading via the 'pip install --upgrade pip' command.
```

In [4]:

```
from __future__ import absolute_import, division, print_function, unicode_literals

import tensorflow as tf
from tensorflow import keras

import numpy as np

print(tf.__version__)
```

WARNING:tensorflow:

TensorFlow's `tf-nightly` package will soon be updated to TensorFlow 2.0.

Please upgrade your code to TensorFlow 2.0:

* https://www.tensorflow.org/beta/guide/migration_guide

Or install the latest stable TensorFlow 1.X release:

* `pip install -U "tensorflow==1.*"`

Otherwise your code may be broken by the change.

1.15.0-dev20190821

Download the IMDB dataset

The IMDB dataset comes packaged with TensorFlow. It has already been preprocessed such that the reviews (sequences of words) have been converted to sequences of integers, where each integer represents a specific word in a dictionary.

The following code downloads the IMDB dataset to your machine (or uses a cached copy if you've already downloaded it):

In [5]:

```
imdb = keras.datasets.imdb

(train_data, train_labels), (test_data, test_labels) = imdb.load_data(num_words=10000)
```

Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/imdb.npz>
17465344/17464789 [=====] - 0s 0us/step

The argument `num_words=10000` keeps the top 10,000 most frequently occurring words in the training data. The rare words are discarded to keep the size of the data manageable.

Explore the data

Let's take a moment to understand the format of the data. The dataset comes preprocessed: each example is an array of integers representing the words of the movie review. Each label is an integer value of either 0 or 1, where 0 is a negative review, and 1 is a positive review.

In [6]:

```
print("Training entries: {}, labels: {}".format(len(train_data), len(train_labels)))
```

Training entries: 25000, labels: 25000

The text of reviews have been converted to integers, where each integer represents a specific word in a dictionary. Here's what the first review looks like:

In [7]:

```
print(train_data[0])
```

```
[1, 14, 22, 16, 43, 530, 973, 1622, 1385, 65, 458, 4468, 66, 3941, 4, 173, 36, 256, 5, 25, 100, 43, 838, 11, 2, 50, 670, 2, 9, 35, 480, 284, 5, 150, 4, 172, 112, 167, 2, 336, 385, 39, 4, 172, 4536, 1111, 17, 546, 38, 13, 447, 4, 192, 50, 16, 6, 147, 2025, 19, 14, 22, 4, 1920, 4613, 469, 4, 22, 71, 87, 12, 16, 43, 530, 38, 7, 6, 15, 13, 1247, 4, 22, 17, 515, 17, 12, 16, 626, 18, 2, 5, 62, 386, 12, 8, 316, 8, 106, 5, 4, 2223, 5244, 1, 6, 480, 66, 3785, 33, 4, 130, 12, 16, 38, 619, 5, 25, 124, 51, 36, 135, 48, 25, 1415, 33, 6, 22, 12, 215, 2, 8, 77, 52, 5, 14, 407, 16, 82, 2, 8, 4, 107, 117, 5952, 15, 256, 4, 2, 7, 3766, 5, 723, 36, 71, 43, 530, 47, 6, 26, 400, 317, 46, 7, 4, 2, 1029, 13, 104, 88, 4, 381, 15, 297, 98, 32, 2071, 56, 26, 141, 6, 194, 7486, 1, 8, 4, 226, 22, 21, 134, 476, 26, 480, 5, 144, 30, 5535, 18, 51, 36, 28, 224, 92, 25, 104, 4, 226, 65, 16, 3, 8, 1334, 88, 12, 16, 283, 5, 16, 4472, 113, 103, 32, 15, 16, 5345, 19, 178, 32]
```

Movie reviews may be different lengths. The below code shows the number of words in the first and second reviews. Since inputs to a neural network must be the same length, we'll need to resolve this later.

In [8]:

```
len(train_data[0]), len(train_data[1])
```

Out[8]:

```
(218, 189)
```

Convert the integers back to words

It may be useful to know how to convert integers back to text. Here, we'll create a helper function to query a dictionary object that contains the integer to string mapping:

In [9]:

```
# A dictionary mapping words to an integer index
word_index = imdb.get_word_index()

# The first indices are reserved
word_index = {k:(v+3) for k,v in word_index.items()}
word_index["<PAD>"] = 0
word_index["<START>"] = 1
word_index["<UNK>"] = 2 # unknown
word_index["<UNUSED>"] = 3

reverse_word_index = dict([(value, key) for (key, value) in word_index.items()])

def decode_review(text):
    return ' '.join([reverse_word_index.get(i, '?') for i in text])
```

```
Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/imdb_word_index.json
1646592/1641221 [=====] - 0s 0us/step
```

Now we can use the `decode_review` function to display the text for the first review:

In [10]:

```
decode_review(train_data[0])
```

Out[10]:

```
"<START> this film was just brilliant casting location scenery story direction everyone's really suited the part they played and you could just imagine being there robert <UNK> is an amazing actor and now the same being director <UNK> father came from the same scottish island as myself so i loved the fact there was a real connection with this film the witty remarks throughout the film were great it was just brilliant so much that i bought the film as soon as it was released for <UNK> and would recommend it to everyone to watch and the fly fishing was amazing really cried at the end it was so sad and you know what they say if you cry at a film it must have been good and this definitely was also <UNK> to the two little boy's that played the <UNK> of norman and paul they were just brilliant children are often left out of the <UNK> list i think because the stars that play them all grown up are such a big profile for the whole film but these children are amazing and should be praised for what they have done don't you think the whole story was so lovely because it was true and was someone's life after all that was shared with us all"
```

Prepare the data

The reviews—the arrays of integers—must be converted to tensors before fed into the neural network. This conversion can be done a couple of ways:

- Convert the arrays into vectors of 0s and 1s indicating word occurrence, similar to a one-hot encoding. For example, the sequence [3, 5] would become a 10,000-dimensional vector that is all zeros except for indices 3 and 5, which are ones. Then, make this the first layer in our network—a Dense layer—that can handle floating point vector data. This approach is memory intensive, though, requiring a `num_words * num_reviews` size matrix.
- Alternatively, we can pad the arrays so they all have the same length, then create an integer tensor of shape `max_length * num_reviews`. We can use an embedding layer capable of handling this shape as the first layer in our network.

In this tutorial, we will use the second approach.

Since the movie reviews must be the same length, we will use the [pad_sequences](https://keras.io/preprocessing/sequence/#pad_sequences) (https://keras.io/preprocessing/sequence/#pad_sequences) function to standardize the lengths:

In [11]:

```
train_data = keras.preprocessing.sequence.pad_sequences(train_data,
                                                         value=word_index["<PAD>"],
                                                         padding='post',
                                                         maxlen=256)

test_data = keras.preprocessing.sequence.pad_sequences(test_data,
                                                         value=word_index["<PAD>"],
                                                         padding='post',
                                                         maxlen=256)
```

Let's look at the length of the examples now:

In [12]:

```
len(train_data[0]), len(train_data[1])
```

Out[12]:

(256, 256)

And inspect the (now padded) first review:

In [13]:

```
print(train_data[0])
```

```
[  1  14  22  16  43 530 973 1622 1385  65 458 4468  66 3941
   4 173  36 256   5  25 100  43 838 112  50 670   2   9
  35 480 284   5 150   4 172 112 167   2 336 385  39   4
 172 4536 1111  17 546  38  13 447   4 192  50  16   6 147
2025  19  14  22   4 1920 4613 469   4  22  71  87  12  16
  43 530  38  76  15  13 1247   4  22  17 515  17  12  16
 626  18   2   5  62 386  12   8 316   8 106   5   4 2223
5244  16 480  66 3785  33   4 130  12  16  38 619   5  25
 124  51  36 135  48  25 1415  33   6  22  12 215  28  77
  52   5  14 407  16  82   2   8   4 107 117 5952  15 256
   4   2   7 3766   5 723  36  71  43 530 476  26 400 317
  46   7   4   2 1029  13 104  88   4 381  15 297  98  32
2071  56  26 141   6 194 7486  18   4 226  22  21 134 476
  26 480   5 144  30 5535  18  51  36  28 224  92  25 104
   4 226  65  16  38 1334  88  12  16 283   5  16 4472 113
103  32  15  16 5345  19 178  32   0   0   0   0   0   0
   0   0   0   0   0   0   0   0   0   0   0   0   0   0
   0   0   0   0   0   0   0   0   0   0   0   0   0   0
   0   0   0   0]
```

Build the model

The neural network is created by stacking layers—this requires two main architectural decisions:

- How many layers to use in the model?
- How many *hidden units* to use for each layer?

In this example, the input data consists of an array of word-indices. The labels to predict are either 0 or 1. Let's build a model for this problem:

In [14]:

```
# input shape is the vocabulary count used for the movie reviews (10,000 words)
vocab_size = 10000

model = keras.Sequential()
model.add(keras.layers.Embedding(vocab_size, 16))
model.add(keras.layers.GlobalAveragePooling1D())
model.add(keras.layers.Dense(16, activation=tf.nn.relu))
model.add(keras.layers.Dense(1, activation=tf.nn.sigmoid))

model.summary()
```

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/tensorflow_core/python/keras/initializers.py:119: calling RandomUniform.__init__ (from tensorflow.python.ops.init_ops) with dtype is deprecated and will be removed in a future version.

Instructions for updating:

Call initializer instance with the dtype argument instead of passing it to the constructor

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/tensorflow_core/python/ops/resource_variable_ops.py:1630: calling BaseResourceVariable.__init__ (from tensorflow.python.ops.resource_variable_ops) with constraint is deprecated and will be removed in a future version.

Instructions for updating:

If using Keras pass *_constraint arguments to layers.

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
embedding (Embedding)	(None, None, 16)	160000

global_average_pooling1d (Gl	(None, 16)	0

dense (Dense)	(None, 16)	272

dense_1 (Dense)	(None, 1)	17
=====		
Total params: 160,289		
Trainable params: 160,289		
Non-trainable params: 0		

The layers are stacked sequentially to build the classifier:

1. The first layer is an `Embedding` layer. This layer takes the integer-encoded vocabulary and looks up the embedding vector for each word-index. These vectors are learned as the model trains. The vectors add a dimension to the output array. The resulting dimensions are: `(batch, sequence, embedding)` .
2. Next, a `GlobalAveragePooling1D` layer returns a fixed-length output vector for each example by averaging over the sequence dimension. This allows the model to handle input of variable length, in the simplest way possible.
3. This fixed-length output vector is piped through a fully-connected (`Dense`) layer with 16 hidden units.
4. The last layer is densely connected with a single output node. Using the `sigmoid` activation function, this value is a float between 0 and 1, representing a probability, or confidence level.

Hidden units

The above model has two intermediate or "hidden" layers, between the input and output. The number of outputs (units, nodes, or neurons) is the dimension of the representational space for the layer. In other words, the amount of freedom the network is allowed when learning an internal representation.

If a model has more hidden units (a higher-dimensional representation space), and/or more layers, then the network can learn more complex representations. However, it makes the network more computationally expensive and may lead to learning unwanted patterns—patterns that improve performance on training data but not on the test data. This is called *overfitting*, and we'll explore it later.

Loss function and optimizer

A model needs a loss function and an optimizer for training. Since this is a binary classification problem and the model outputs a probability (a single-unit layer with a sigmoid activation), we'll use the `binary_crossentropy` loss function.

This isn't the only choice for a loss function, you could, for instance, choose `mean_squared_error` . But, generally, `binary_crossentropy` is better for dealing with probabilities—it measures the "distance" between probability distributions, or in our case, between the ground-truth distribution and the predictions.

Later, when we are exploring regression problems (say, to predict the price of a house), we will see how to use another loss function called mean squared error.

Now, configure the model to use an optimizer and a loss function:

In [15]:

```
model.compile(optimizer='adam',
              loss='binary_crossentropy',
              metrics=['acc'])
```

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/tensorflow_core/python/ops/nn_impl.py:183: where (from tensorflow.python.ops.array_ops) is deprecated and will be removed in a future version.
Instructions for updating:
Use tf.where in 2.0, which has the same broadcast rule as np.where

Create a validation set

When training, we want to check the accuracy of the model on data it hasn't seen before. Create a *validation set* by setting apart 10,000 examples from the original training data. (Why not use the testing set now? Our goal is to develop and tune our model using only the training data, then use the test data just once to evaluate our accuracy).

In [16]:

```
x_val = train_data[:10000]
partial_x_train = train_data[10000:]

y_val = train_labels[:10000]
partial_y_train = train_labels[10000:]
```

Train the model

Train the model for 40 epochs in mini-batches of 512 samples. This is 40 iterations over all samples in the `x_train` and `y_train` tensors. While training, monitor the model's loss and accuracy on the 10,000 samples from the validation set:

In [17]:

```
history = model.fit(partial_x_train,  
                    partial_y_train,  
                    epochs=40,  
                    batch_size=512,  
                    validation_data=(x_val, y_val),  
                    verbose=1)
```

Train on 15000 samples, validate on 10000 samples

Epoch 1/40

15000/15000 [=====] - 1s 34us/sample - loss: 0.6920 - acc: 0.6093 - val_loss: 0.689

9 - val_acc: 0.7129

Epoch 2/40

15000/15000 [=====] - 0s 28us/sample - loss: 0.6863 - acc: 0.7362 - val_loss: 0.682

2 - val_acc: 0.7393

Epoch 3/40

15000/15000 [=====] - 0s 28us/sample - loss: 0.6745 - acc: 0.7486 - val_loss: 0.667

5 - val_acc: 0.7376

Epoch 4/40

15000/15000 [=====] - 0s 27us/sample - loss: 0.6543 - acc: 0.7611 - val_loss: 0.644

2 - val_acc: 0.7365

Epoch 5/40

15000/15000 [=====] - 0s 28us/sample - loss: 0.6242 - acc: 0.7841 - val_loss: 0.612

4 - val_acc: 0.7715

Epoch 6/40

15000/15000 [=====] - 0s 27us/sample - loss: 0.5856 - acc: 0.8062 - val_loss: 0.574

6 - val_acc: 0.8023

Epoch 7/40

15000/15000 [=====] - 0s 27us/sample - loss: 0.5413 - acc: 0.8252 - val_loss: 0.532

6 - val_acc: 0.8154

Epoch 8/40

15000/15000 [=====] - 0s 28us/sample - loss: 0.4957 - acc: 0.8429 - val_loss: 0.492

5 - val_acc: 0.8284

Epoch 9/40

15000/15000 [=====] - 0s 28us/sample - loss: 0.4522 - acc: 0.8568 - val_loss: 0.455

4 - val_acc: 0.8407

Epoch 10/40

15000/15000 [=====] - 0s 28us/sample - loss: 0.4129 - acc: 0.8689 - val_loss: 0.423

9 - val_acc: 0.8483

Epoch 11/40

15000/15000 [=====] - 0s 27us/sample - loss: 0.3791 - acc: 0.8781 - val_loss: 0.397

3 - val_acc: 0.8558

Epoch 12/40

15000/15000 [=====] - 0s 27us/sample - loss: 0.3502 - acc: 0.8865 - val_loss: 0.376

6 - val_acc: 0.8586

Epoch 13/40

15000/15000 [=====] - 0s 28us/sample - loss: 0.3266 - acc: 0.8920 - val_loss: 0.358

2 - val_acc: 0.8652

Epoch 14/40

15000/15000 [=====] - 0s 27us/sample - loss: 0.3052 - acc: 0.8977 - val_loss: 0.344

6 - val_acc: 0.8691

Epoch 15/40

15000/15000 [=====] - 0s 27us/sample - loss: 0.2875 - acc: 0.9019 - val_loss: 0.333

3 - val_acc: 0.8721

Epoch 16/40

```
15000/15000 [=====] - 0s 27us/sample - loss: 0.2716 - acc: 0.9077 - val_loss: 0.323
8 - val_acc: 0.8738
Epoch 17/40
15000/15000 [=====] - 0s 27us/sample - loss: 0.2571 - acc: 0.9119 - val_loss: 0.315
9 - val_acc: 0.8764
Epoch 18/40
15000/15000 [=====] - 0s 27us/sample - loss: 0.2442 - acc: 0.9161 - val_loss: 0.309
1 - val_acc: 0.8790
Epoch 19/40
15000/15000 [=====] - 0s 28us/sample - loss: 0.2327 - acc: 0.9194 - val_loss: 0.303
3 - val_acc: 0.8804
Epoch 20/40
15000/15000 [=====] - 0s 27us/sample - loss: 0.2222 - acc: 0.9234 - val_loss: 0.299
3 - val_acc: 0.8800
Epoch 21/40
15000/15000 [=====] - 0s 26us/sample - loss: 0.2118 - acc: 0.9287 - val_loss: 0.295
7 - val_acc: 0.8812
Epoch 22/40
15000/15000 [=====] - 0s 27us/sample - loss: 0.2029 - acc: 0.9299 - val_loss: 0.292
6 - val_acc: 0.8833
Epoch 23/40
15000/15000 [=====] - 0s 27us/sample - loss: 0.1940 - acc: 0.9339 - val_loss: 0.290
9 - val_acc: 0.8833
Epoch 24/40
15000/15000 [=====] - 0s 27us/sample - loss: 0.1863 - acc: 0.9383 - val_loss: 0.288
5 - val_acc: 0.8837
Epoch 25/40
15000/15000 [=====] - 0s 27us/sample - loss: 0.1784 - acc: 0.9427 - val_loss: 0.286
7 - val_acc: 0.8841
Epoch 26/40
15000/15000 [=====] - 0s 27us/sample - loss: 0.1714 - acc: 0.9451 - val_loss: 0.286
6 - val_acc: 0.8838
Epoch 27/40
15000/15000 [=====] - 0s 26us/sample - loss: 0.1648 - acc: 0.9479 - val_loss: 0.285
8 - val_acc: 0.8852
Epoch 28/40
15000/15000 [=====] - 0s 27us/sample - loss: 0.1584 - acc: 0.9515 - val_loss: 0.285
4 - val_acc: 0.8849
Epoch 29/40
15000/15000 [=====] - 0s 27us/sample - loss: 0.1529 - acc: 0.9531 - val_loss: 0.286
8 - val_acc: 0.8836
Epoch 30/40
15000/15000 [=====] - 0s 28us/sample - loss: 0.1471 - acc: 0.9546 - val_loss: 0.285
7 - val_acc: 0.8853
Epoch 31/40
15000/15000 [=====] - 0s 27us/sample - loss: 0.1413 - acc: 0.9579 - val_loss: 0.286
2 - val_acc: 0.8861
Epoch 32/40
```



```

15000/15000 [=====] - 0s 27us/sample - loss: 0.1359 - acc: 0.9600 - val_loss: 0.287
3 - val_acc: 0.8861
Epoch 33/40
15000/15000 [=====] - 0s 27us/sample - loss: 0.1308 - acc: 0.9620 - val_loss: 0.289
1 - val_acc: 0.8851
Epoch 34/40
15000/15000 [=====] - 0s 26us/sample - loss: 0.1263 - acc: 0.9637 - val_loss: 0.290
3 - val_acc: 0.8863
Epoch 35/40
15000/15000 [=====] - 0s 27us/sample - loss: 0.1221 - acc: 0.9643 - val_loss: 0.292
5 - val_acc: 0.8857
Epoch 36/40
15000/15000 [=====] - 0s 27us/sample - loss: 0.1175 - acc: 0.9672 - val_loss: 0.294
0 - val_acc: 0.8853
Epoch 37/40
15000/15000 [=====] - 0s 27us/sample - loss: 0.1132 - acc: 0.9688 - val_loss: 0.296
0 - val_acc: 0.8847
Epoch 38/40
15000/15000 [=====] - 0s 28us/sample - loss: 0.1091 - acc: 0.9699 - val_loss: 0.299
0 - val_acc: 0.8837
Epoch 39/40
15000/15000 [=====] - 0s 27us/sample - loss: 0.1059 - acc: 0.9710 - val_loss: 0.301
6 - val_acc: 0.8827
Epoch 40/40
15000/15000 [=====] - 0s 27us/sample - loss: 0.1018 - acc: 0.9727 - val_loss: 0.303
7 - val_acc: 0.8831

```

Evaluate the model

And let's see how the model performs. Two values will be returned. Loss (a number which represents our error, lower values are better), and accuracy.

In [18]:

```

results = model.evaluate(test_data, test_labels)

print(results)

```

```

25000/25000 [=====] - 1s 20us/sample - loss: 0.3234 - acc: 0.8728
[0.3234443087768555, 0.87276]

```

This fairly naive approach achieves an accuracy of about 87%. With more advanced approaches, the model should get closer to 95%.

Create a graph of accuracy and loss over time

`model.fit()` returns a `History` object that contains a dictionary with everything that happened during training:

In [19]:

```
history_dict = history.history
history_dict.keys()
```

Out[19]:

```
dict_keys(['loss', 'acc', 'val_loss', 'val_acc'])
```

There are four entries: one for each monitored metric during training and validation. We can use these to plot the training and validation loss for comparison, as well as the training and validation accuracy:

In [20]:

```
import matplotlib.pyplot as plt

acc = history_dict['acc']
val_acc = history_dict['val_acc']
loss = history_dict['loss']
val_loss = history_dict['val_loss']

epochs = range(1, len(acc) + 1)

# "bo" is for "blue dot"
plt.plot(epochs, loss, 'bo', label='Training loss')
# b is for "solid blue line"
plt.plot(epochs, val_loss, 'b', label='Validation loss')
plt.title('Training and validation loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()

plt.show()
```


<Figure size 640x480 with 1 Axes>

In [21]:

```
plt.clf() # clear figure

plt.plot(epochs, acc, 'bo', label='Training acc')
plt.plot(epochs, val_acc, 'b', label='Validation acc')
plt.title('Training and validation accuracy')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend()

plt.show()
```



In this plot, the dots represent the training loss and accuracy, and the solid lines are the validation loss and accuracy.

Notice the training loss *decreases* with each epoch and the training accuracy *increases* with each epoch. This is expected when using a gradient descent optimization—it should minimize the desired quantity on every iteration.

This isn't the case for the validation loss and accuracy—they seem to peak after about twenty epochs. This is an example of overfitting: the model performs better on the training data than it does on data it has never seen before. After this point, the model over-optimizes and learns representations *specific* to the training data that do not *generalize* to test data.

For this particular case, we could prevent overfitting by simply stopping the training after twenty or so epochs. Later, you'll see how to do this automatically with a callback.

In []:

In []: